

Deformable Muscle Models for Motion Simulation

Tomáš Janák and Josef Kohout

*Department of Computer Science and Engineering, University of West Bohemia,
Univerzitní 8, 30100, Plzeň, Czech Republic*

*NTIS – New Technologies for the Information Society, University of West Bohemia,
Univerzitní 22, 30100, Plzeň, Czech Republic*

Keywords: Musculoskeletal Model, Medical Simulation, Soft-Body, Deformable Objects, Collision Detection, Mass-Spring System.

Abstract: This paper presents a methodology for interactive muscle simulation. The fibres of individual muscles are represented by particles connected by springs, thus creating a deformable model of the muscle. In order to be able to describe human musculoskeletal system, contact between pairs of muscles as well as muscles and bones must be accounted for. Therefore, collision detection and response mechanism which allows both types of contact (soft body vs. rigid body and soft vs. soft body) is presented. The solution is a part of a project dedicated to improvement of the effectiveness of osteoporosis prediction and treatment.

1 INTRODUCTION

The musculoskeletal system of modern human is often a subject to various medical conditions, from minor aches to serious diseases such as osteoporosis. To improve the efficiency of treatment, virtual models can be used to simulate motion of the patient during some situations (walking, jumping etc.) in order to better understand how such workload affects muscles and bones in these situations. The creation of the whole musculoskeletal model that could be tuned to a specific patient is a complicated procedure with many steps that are outside the scope of this paper. The focus here is solely on how to model muscles during the simulated motion, assuming their initial geometry is given.

In terms of computer simulation, a muscle is a typical example of a soft-body object, i.e. its shape is elastically changing in response to external forces. There are two general approaches for simulating the musculoskeletal system in motion. The first, simpler, approach defines the animation by movement of the bones and then deforms the muscles according to the interaction with bones or potential obstacles. That means that the muscles are deformed as a result of the animation instead of being the initiators of the animation. The second approach models individual muscles as they actually work in real world, i.e. they are the initiators of forces that move the bones. The source of the animation then are the forces acting on

the bones produced by individual muscles. Although the second approach is possible to use (Lee et al, 2009), it is obviously computationally much more demanding and will not be considered here.

As the simulated patient moves through the virtual scene, the muscles interact with the bones, possible obstacles in the scenes and also each other. Therefore, apart from the obvious need to be able to update the geometry in every step of the simulation, the crucial part of the muscle simulation is efficient collision detection, which identifies parts of the muscle geometry that should be updated. There are many different problems in computer graphics and related fields which rely on collision detection. However, in most cases the situation is slightly less complicated and satisfied with rigid object vs. rigid object, or one rigid vs. one soft-body object collision detection. In the case of musculoskeletal system simulation, all possible combinations of collision situations happen at the same time. Furthermore, each muscle is modelled separately. Therefore, the neighbouring muscles are in fact in a state of continuous collision, affecting each other in every single step of the motion. That results in the need for a robust collision handling methodology.

One could argue that all the adjacent muscles could be treated as one large lump of tissue, but that would limit the options the solution provides. First, the operator will often be interested in analysis of only a single (or several) muscles and lumping all the muscles together would complicate that. Also,

even non-adjacent muscles can touch each other during various motions (e.g. thighs and calves touching while kneeling etc.) and therefore the collisions would have to be checked for and treated nevertheless. Lastly, connecting individual muscles together does not comply with the physical reality. Even when attached to the same bone, muscles can slide over each other and change their relative positions by a significant margin. This would be difficult to simulate if the muscles were as one.

The idea behind this work was to create a framework that would enable medical operators to quickly create an interactive simulation of motion, i.e. ideally a real-time simulation, but at least a frame every few seconds. Such simulation would then be used for a coarse assessment of the situation at hand and only after finding out the most critical points during the motion, a more accurate (but also time consuming) algorithm would be used for a precise evaluation of those few points. Hence, even though the simulation has to be realistic in order to be useful, it does not aim to be perfect. The framework was created as a part of the EU-funded VPHOP project (www.vphop.eu), which is aimed on developing technologies for better prediction of bone fracture risks in order to be able to provide more effective treatment of osteoporosis.

On the following pages, the paper will present the created framework. There is no groundbreaking new algorithm presented, rather the contribution is in assessment of existing algorithms of soft-body simulation and collision detection and “tweaking” those for the purposes of the particular problem described above.

This paper is structured as follows. After a brief summary of the previous work done in related fields in Section 2, the main part of the paper will follow with detailed description and some implementation details of the used muscle model and collision handling mechanism (Section 3). Section 4 will conclude with experimental measurements of the framework's performance.

2 STATE OF THE ART

2.1 Soft-Body Models

A soft-body represent a stiff, but deformable object. The shape of the object changes according to external forces, but at the same time the object resists those forces and tries to maintain its original (“rest”) shape. In general, soft-body models can be classified as either heuristic or continuum mechanics, depend-

ing on whether the model behaves according to actual elasticity principles or some heuristic that aims only to produce a plausible, although not physically accurate, visualisation. A continuum mechanics approach offers better fidelity and also more user comfort because all the parameters of the model are actual measurable physical properties. Heuristic approaches usually require a lot of experimenting to find suitable parameters, but they are significantly faster during the run-time of the simulation. The most common method for the continuum mechanics approach is the Finite Element Method (FEM), while probably the most common heuristic used for soft-bodies are the mass-spring systems (MSS), which were also chosen for this work in order to be able to achieve faster execution.

As the name implies, the object (its surface only or the whole volume) modelled by MSS is discretized into a set of point masses (particles) which are connected by springs. A particle is defined by mass and position. A spring is defined by stiffness and damping coefficients, rest length (length of the spring in rest position) and the two particles it connects. Hooke's law describes the force acting on the connected particles by a “spring equation”, while the movement of each particle is described by usual Newtonian mechanics. All the particles and spring parameters do not have any connection to the object they represent, so it can be difficult to set these parameters to express the behaviour of the simulated material correctly.

The MSS methods are popular especially in the field of cloth modelling and a lot of materials on this topic can be found for example in a recent survey by Long et al. (2011). The application of MSS for medical purposes is much less common, mainly due to the lack of physical accuracy. Nedel and Thalmann (1998) were one of the first to exploit MSS for muscle simulation. They model only the surface of the muscle, aligning vertices of the surface triangular model with the particles of the MSS and its edges with the springs. They introduce additional angular springs to limit torsion and also volume loss of the muscle. This model is used for visualization. External forces yield from the underlying action line model that approximates all the muscle fibres of a muscle by one poly-line. While fast to process, an assumption that all the muscle fibres have the same length can lead to a wrong force load analysis.

Villard et al. (2008) combined the elastic springs of MSS with solid parts to create a system which is still elastic, yet compressible only to a certain extent and successfully used it to model diaphragm motion. Hui and Tang (2009) used the MSS to model tendon

motion and deformation. Their system was also based on the work of Nedel and Thalmann (1998), using additional flexion and angular springs in order to contain the torsion and other unwanted deformations.

A recent survey of Lee et al. (2012) compiles a comprehensive set of various approaches to muscle modelling, including FEM and MSS based models and also fast data driven approaches suitable for non-medical purposes. A reader will find many references to noteworthy publications there.

2.2 Collision Detection

Detecting collisions of two complex objects in fact means detecting collisions between two groups of primitives, i.e. triangles in the most common case of triangular meshes. As there can be thousands or even millions of triangles involved in the scene, almost all collision detection (CD) algorithms include some pruning phase that limits the number of primitives that have to be checked in the phase of piecewise tests.

Methods based on Bounding Volume Hierarchies (BVH) divide the object into a hierarchical structure of simple wrapping geometric shapes such as axis aligned (AABB) or oriented (OBB) bounding boxes, spheres etc. CD of two objects starts with testing the root node of the BVH of one object, which bounds the whole object, with the root of the other. If an overlap is detected, the following levels of the hierarchy are traversed and tested until there is no intersection detected or until the lowest level is reached. In the latter case, the primitives stored in the leaf nodes are subjected to piecewise testing.

A fundamental flaw in the BVH approach in the context of soft-bodies is that as the object changes shape, the bounding structure constructed above it may become invalid. Therefore, before each CD, the BVH needs to be validated and updated if needed. There are two common update mechanism – refitting and rebuilding. Refitting only updates individual bounding volumes, usually inflating them in order to accommodate primitives that moved outside their bounds. When the changes in the shape are too vast for refitting, the BVH needs to be rebuilt. Levels in the hierarchical structure are removed bottom-up until lowest valid level is reached and then the hierarchy is built again. These operations obviously need to be fast, which is why the simplest bounding primitives – usually AABBs – are most commonly used, even though they might not have such a tight fit (and therefore less “false positive” overlap tests) as other shapes used in CD for rigid objects.

Larsson and Akenine-Möller (2006) proposed a robust BVH solution for CD of deformable and even breakable objects. Their solution exploits assumed temporal coherence of subsequent CD steps – the nodes, that were used in previous CD query are marked as active and refitted during the update step. Also, the nodes are validated in that step by comparing the volume of a node to sum of volumes of its children. If the difference is too large, the children nodes are destroyed. However, they are not rebuilt until they are needed.

Mendoza and Sullivan (2006) introduced interruptible algorithm using BVH for time critical CD between soft-bodies. Tang et al. (2009) created a parallel BVH-based algorithm for continuous CD. They introduced several advanced pruning concepts that allowed them to achieve interactive frame rates for scenes with several tens of thousands triangles, showing that BVH can be used efficiently for soft-body CD.

Other methods suitable for the soft-body vs. soft-body CD include spatial hashing techniques (Teschner et al, 2003), (Hoppe and Lefebvre, 2006) or methods using the layered depth image (Faure et al, 2008). As the BVH methods were chosen for the solution, details about these methods will not be discussed here.

3 SOLUTION DESIGN

This chapter will describe the devised solution for muscle simulation. The assignment can be defined as follows. First, the solution must correctly deform each muscle in a given group of muscles based on the interaction with their surroundings (including other muscles) while undergoing a predefined motion animation. Interpenetration of individual objects must be prevented and the results must be presented in an interactive rate for a single area of interest (e.g. one lower limb or one upper limb etc.).

3.1 Overall Pipeline

The input data consist of triangular meshes of the muscles and bones, polylines representing muscle fibres and transformation matrices for the bones for each time step, i.e. the motion animation. The user interface was designed in such a way that enables rendering an arbitrary point on the timeline (time point), i.e. it is not assumed that the animation will be temporarily continuous. For this reason, whenever a given time point is processed, the rest pose of the muscles and bones is the starting point – all




Figure 1: A detailed visualization of muscle fibres in the rest position (left), after rigid transformation to the current position (middle) and after stabilization of the MSS (right).



Figure 2: Overview of the bone positions for the situation depicted on Figure 1. Left is the rest position (related to Figure 1-left) and right is the position during walking (related to Figure 1-middle and right).

transformations are related to this position.

The first step is pre-processing, which converts the muscle fibre polylines into particles and springs. This is controlled by a single parameter, the number of particles per fibre N . The fibres are divided into $N-1$ uniform intervals, but instead of placing the particles as the end points of those intervals, their position is randomized within each interval. This is done because when N is small and the fibres of the muscle are roughly the same length, which is often the case, the particles placed without randomization tend to create clusters with large gaps in between them and that negatively affects the CD precision (the CD uses the particles as an approximation of the muscle surface, therefore they should be scattered

along the surface as much as possible, details in section 3.3).

Once the particles are created, the spring interconnections between them are generated using a user-selected pattern (see 3.2 for list of implemented patterns) and then all this is stored as simple arrays, floating point for the particle positions and pair of integer indices for each spring. Some other minor structures are created during the pre-processing as well, such as associations of particles with bones closest to them, markings of boundary particles etc. (see section 3.2 for details). This representation is stored and reused through different time points unless the user changes some key parameters.

To process a given time point of the motion, the transformation matrices are first applied to the bones. The muscles (the triangular surfaces as well as the particles) are transformed as well, using an interpolated transformation matrix of the nearest bones. After this basic rigid transformation, the soft-body simulation takes place.

The simulation is an iterative process consisting of two major parts - MSS update followed by collision handling. During the MSS update, new positions of particles are computed based on the forces acting in the MSS. These new positions will likely introduce some collisions with neighbouring objects and handling those collisions in return adds some new forces into the system. In an ideal situation, after iterating this loop a finite number of times, the MSS should reach an equilibrium state where no particles move anymore. Then the final shape of the muscles is ready and can be displayed.

However, an absolute equilibrium is obviously practically impossible to achieve during computer simulation. There are three possibilities of how to end the simulation loop, each equally simple to im-

plement. In order to achieve best fidelity of the output, an average displacement of the particles can be monitored and when it falls beneath a given threshold, the state is considered final. Or, if the rendering time is more critical, a given time window, e.g. based on the desired frame rate, can be assigned to the computation and the output is produced immediately after this window is depleted. A compromise is to set a fixed number of iterations.

After the final particle positions – which define the final shape – are computed, the surface model is updated (see 3.2) and visualized. Figures 1 and 2 show the main stages of the process on an example of a right leg performing a common walking step.

As there are generally many time points in the animation, the created result is discarded after the user moves onto another time point. This unfortunately slows the execution when a whole continuous animation is wanted, as each time point is handled separately, without exploiting the temporal coherence in any way. For example, the rigid transformations from rest position to the current position could be omitted altogether. Moreover, as it can be expected that there will be only relatively small changes in consecutive frames, a much smaller number of iterations of the simulation loop would be required to get to a plausible state in consecutive time points. Nevertheless, the application for which this method was designed required arbitrary time point changes. Should the requirements change however, the algorithm would not be difficult to modify.

3.2 Muscle Model

There are two muscle models – the triangular surface and the other is the muscle fibre model. As was mentioned before, the key model used in the workload analysis is the fibre one, the surface model is used only for the visualization. MSS are used to represent the muscles. The MSS solver was already implemented in the target framework, more details about it can be found in Zelený (2011).

The solver is computing particle positions in every time step of the simulation loop. As the particles interconnected by springs along the muscle fibre actually represent this fibre, it is straightforward to obtain its new shape after the simulation ends since it is still the same polyline, only its vertices have different positions. However, it is not as easy to obtain the new shape of the muscle surface. For this reason, some relation between the particles and the triangular model has to be established. We use mean value coordinates (MVC) for this purpose, described

by Ju et al. (2005) as follows:

The mean value interpolation interpolates a given function $f(x)$ defined on a closed surface by a function $g(v)$, $v \in \mathbb{R}^3$ by projecting a point $p(x)$ of the closed surface on a unit sphere centered at v . Then the function value associated with $p(x)$ is weighted by $w = [p(x) - v]^{-1}$ and integrated over the sphere. To ensure affine invariance, the result is divided by the weight function integrated over the sphere S . Equation (1) shows the result:

$$g(v) = \frac{\int w(x, v) f(x) dS}{\int w(x, v) dS} \quad (1)$$

The authors then continue to derive the MVC for closed polygons and mainly triangular meshes. Equation (2) computes the weights for point v in regards to a given triangle with vertices p_i , $i = 1, 2, 3$. n_i are normal vectors of the three triangles $vq_i q_{i+1}$ where q_i are the vertices of a spherical triangle constructed by projecting the original triangle $p_1 p_2 p_3$ on the unit sphere. m is the “mean vector” which is an integral of the outward normal vector over the spherical triangle surface (which is not difficult to evaluate using a few trigonometric functions, see Ju et al. (2005) for details). Weights are computed this way for all triangles and summed analogically as in equation (1) (the integrals are replaced by sums over all the primitives) to obtain the MVC.

$$w_i = \frac{n_i \cdot m}{n_i \cdot (p_i - v)} \quad (2)$$

In our case, the boundary (located on outmost fibres) particles are triangulated and used as the source mesh, i.e. each particle is the $p(x)$ acting in equations (1) and (2). Each vertex of the muscle surface mesh is then treated as the target point (v in (1) and (2)) and its MVC are computed and stored as a pre-processing. After the particles move during the soft-body simulation, shape of the surface model is updated simply by recalculating position of each vertex as a linear combination of its MVC and the current particle positions. Figure 3 shows the Gluteus maximus and Iliacus muscles, which change their shapes significantly during the simulated motion – in this case simple walking. The MVCs are used to deform the surface from the rest pose to the final pose while conserving the smoothness of the surface.

To simulate the tendon attachment of each muscle, the particles on ends of the fibres are set as fixed. The position of fixed particles is not updated in the soft-body simulation, only during the rigid transformations.




Figure 3: Surfaces of the Gluteus maximus and Iliacus muscles. Up in rest pose (standing), down in the final pose (walking).

There are many ways, or templates, of how to generate the springs connecting the particles. The more springs there are, the longer the computation will be, but on the other hand, systems with low spring count tend to converge to the equilibrium state slower. Several configurations were tested: cubic lattice with 6 or 26 springs per particle (the particles are treated as vertices of a cubic grid with either only the sides of the cubes being the springs, or all sides and both face and space diagonals being the springs); delaunay tetrahedralization (springs are generated as edges of tetrahedrons in a delaunay

tetrahedralization of the particles); and N nearest neighbours (each particle is connected with a user defined number of closest particles). The impact of the choice of the layout is discussed in Section 4.1.

3.3 Collision Handling

Collision handling is responsible for ensuring that no objects penetrate each other. However, as section 3.2 established, the surfaces of the muscles are not involved in the simulation at all. That means that one would have to update the surface mesh in each step of the simulation loop, detect collisions, propagate the response onto the particles via the MVC and continue with the simulation. This would significantly increase the computation time. Another problem is that the surfaces of individual muscles can actually intersect even in the initial position. That is not an error in the data – some muscles can, and do, interweave each other, which is easily simulated by using the fibre model, but difficult when using the encapsulating surface.

To bypass these problems, the particle model itself is used for the collision detection (CD) instead. Obviously, the particles represent point masses and therefore do not have any volume. Also, they do not necessarily trace the surface of the muscle. To change that, the particles are thought of as spheres with a given non-zero radius. Interpenetrations between these sets of spheres for different muscles can then be detected and removed. At the same time their position coincides with the particle positions and therefore the collision response consists of nothing more than updating the positions of the particles. Another pleasing by-product of this design is simplification of the CD itself, as detecting and resolving collisions of set of spheres is generally simpler than in the case of triangle meshes – sphere vs. sphere collision test is a simple comparison of their distance and the sum of their radii. Also, only the particles that are on the “boundary” fibres, i.e. fibres closest to the surface, have to be accounted for – the internal particles will not collide with other objects due to the internal forces of the MSS.

The following heuristic has been used to obtain the radius for each sphere. The closest particle to each vertex of the surface mesh is found. Note that one particle can be the closest one to several vertices. To make a compromise between covering enough of the surface of the muscle and the least of the outside space, the radius of each sphere is set to an average of the distances between the centre of the particle and the associated vertices. Then the CD mechanism is employed to detect particles that inter-

sect and their radii are decreased so that the intersection are removed, which effectively removes some very large particles that could have been generated. To increase the coverage of the surface without increasing the volume excessively, more particles per fibre can be used. However, that comes at a price of higher memory and time demands.

For collision detection between the muscles and bones, the same mechanism is used. The bones are also converted to a set of spheres that approximates their surface and then the muscle vs. bone CD is processed in the same way as muscle vs. muscle. The position of the spheres representing the bone can be arbitrary (unlike in the case of muscles) and therefore the spheres can trace the surface more closely without taking up too much excess space.

The CD itself uses the BVH mechanism based on the method by Larsson and Akenine-Möller (2006), using AABBs, subdivided into even octants (i.e. the division lines always passes through the midpoint of the parent box). This method was chosen over the others mentioned in Section 2.2, due to its universality and good performance documented in the aforementioned work. Simple AABBs are used, because they are faster to update as the object changes shape. The bounding boxes are subdivided into octants in each level of the subdivision. The division lines always pass through the midpoint of the parent box. Whenever new object is added to the scene, the bounding box is constructed for it only on the parent level, i.e. without any subdivision.

The rebuilding proposed by the original authors (see 2.2) is not used. After testing, it has been found out, that the proposed rules for determining whether to rebuild part of the hierarchy does not work well in this case. The CD is on average 5% faster when only refitting is used. That is actually not very surprising, as it can be expected that the shape of the muscles will not change drastically during one simulation step, when the particles only head toward the equilibrium state. Also, the bounding boxes of the bones obviously never need to be updated, because they do not move (not in the span of one simulation step) or change shape.

Each muscle in the scene is then tested against each other muscle and also each bone. The BVHs of the objects are traversed and subdivided only when needed. The maximal number of recursion steps as well as the minimal number of primitive per node are either specified by a user or their default values, empirically determined to be 8 and 10, respectively, are used. Once either of the limits is reached, and there is still a collision detected between the nodes, a pair of lists containing the primitives in the two

colliding nodes is outputted and the piecewise test of the primitives in these lists is done.

When the piecewise test detects a collision, the difference between the distance of the two particles and the sum of their radii is computed. Then each particle is moved in an opposite direction to each other by half of this distance, or eventually, if a fixed particle is involved, the unfixed particle is moved by the whole distance (spheres that represent bones are treated as fixed particles). This way the particles end up just touching each other. To take into account that a given particle can collide with multiple other particles, the displacement vectors are stored and accumulated for each particle and only after all test have been made, their superposition is applied to the initial particle positions.

After all tests are finished, the results are applied (i.e. the accumulated displacement vectors are used to update the positions) and a new iteration may begin. The particles that collided in a given time step are treated as stopped – they are assigned a zero velocity. This ensures that they do not rebound after collision, which is natural for the kind of elastic behaviour that is being simulated.

4 EXPERIMENTS

Our approach was implemented in C++ using the VTK library (<http://www.vtk.org>) and integrated in a MAF based (<http://www.openmaf.org>) application-created for the VPHOP project mentioned earlier.




Figure 4: Overview of the used dataset.

Tests were done on a PC with Intel Core i7-3770 (3,4GHz, 4 physical cores with hyper-threading) CPU, 8GB RAM, MS Windows 7 64-bit OS.

One data set was used for all the tests. It consists of MRI footage of pelvis and legs, fused with a motion capture data of a walking human. A total of twenty three muscles are available for testing in the data set, ten of which are on the pelvis and thirteen on the right thigh. The visualization of this dataset is on Figure 4, with muscles rendered as surfaces.

4.1 Spring Layout

The number of iterations needed to achieve the final shape of the soft-body affects the computational time the most. The spring layout has a significant impact on this number, therefore several layouts were tested. The test processed 1500 iterations of the simulation for each tested method, measuring the average displacement of particles, i.e. the difference in position of each particle between two successive iterations. In an ideal state, the particles would not move at all once the equilibrium position is achieved. However, this is unlikely to happen in the simulation. Rather than zero displacement, an oscillation of the displacement values is a sign of the final state (the positions will oscillate due to the continuous collisions between adjacent particles).

The tested layouts were: Cubic lattice models with 6 and 26 springs per particle, Delaunay tetrahedralization (DT) and 15 nearest neighbours (see section 3.1). The times for those methods were 111.12s, 129.6s, 124.19s and 121.37s respectively, but please note that these times are listed only for comparison between different spring layouts and do not reflect the performance of the final solution – see section 4.3 for a complete analysis of time consumption. The thirteen thigh muscles were used for these tests.

Figure 5 shows the absolute displacements for six of the tested muscles (for better clarity of the chart, some muscles were omitted) for the DT. The displacement rises in the beginning as the particles start to move faster due to the forces introduced into the system by the initial rigid transformations. As they reach the proximity of their final positions, their movement slows down (after around 80 iterations), ideally stopping entirely after more iterations.

The 26-cubic lattice and 15 nearest neighbours have fairly similar progression as the DT on Figure 5, while the 6-cubic lattice displays almost an order lower differences. However, note that the absolute value of displacement measures only the differences between consecutive iterations, not how close are




Figure 5: Absolute displacements for several muscles simulated with the DT spring layout.

these positions to the final state. It is desirable for the differences to actually be as high as possible. That would mean that the model is moving fast towards the final configuration of positions.

Figure 6 depicts the displacements for the Adductor Longus muscle (as Figure 5 documents, the progression is quite similar for all the muscles, therefore the choice of test subject is not important) using four different layouts and including a linear regression of the trend for each layout. Several first iterations were not accounted for, because it is the later “stabilization” phase where the differences between individual layouts have the highest impact.

It is apparent that the 6-cubic lattice layout converges at the slowest rate, while the 26-cubic is the fastest. Even though the 6-cubic layout is the fastest per iterations, much more iterations will be needed in order for the model to reach the equilibrium and therefore the execution will be the slowest. The remaining two methods offer a reasonable compromise. Various count of neighbours in the nearest neighbour method were tested, revealing almost strictly linear relation between the number of springs and convergence speed.

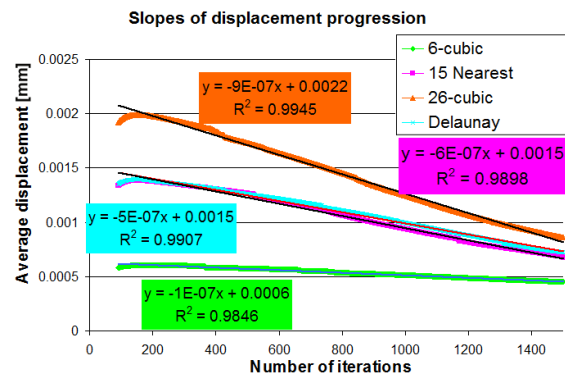


Figure 6: Slopes of particle position displacement for the Adductor Longus muscle.

The 6-cubic lattice is clearly the least suitable option. Although 26-cubic lattice provides fastest convergence, the nearest neighbour or DT layouts also provide fast convergence while having slightly lesser memory and computation time demands. Moreover, if the nearest neighbour is used, the number of neighbours can be provided to the user as an additional parameter to control the model (also note that when using 26 neighbours, the behaviour will be very similar to the 26-cubic). Hence, the nearest neighbour is the suggested solution.

4.2 Deformation Quality

To compare the results of various deformation methods, volume preservation is often used. It is disputable whether this is a relevant metric, since the assumption that a muscle retains its volume does not hold in general. Nevertheless, it is one of few measurable characteristics of different methods and can reveal some relations between them. The error in muscle volume preservation was on average 2.71% for the used dataset. Although larger than the other, more precise method that was available as part of the project solution (which averaged on 0.08%), up to 6% errors are considered to be acceptable and therefore the presented method is usable. Moreover, by using more fibres for the model and a larger stiffness coefficient for the springs, the volume loss should decrease.

More relevant evaluation method is a visual check and comparison with real data by an expert. Because it is almost impossible to obtain real MRI images of the patient in different poses, medical literature must suffice as the data in this case. The implemented solution was handed over to a partner facility (Istituto ortopedico Rizzoli, Italy, <http://www.ior.it>) to do such evaluation test. The key observation was that our fibre model shows “reasonable consistency with the behaviour reported in literature”.

The partner further investigated changes in fibre length during motion. In general, the method performed well, emulating behaviour described in medical literature closely. The only complication arose when the muscle was shortened significantly, as the MSS-driven model tend to buckle as it was resisting the shortening. Lower stiffness coefficients remove this problem, but, as mentioned before, lower stiffness also tend to result in higher volume error. To conclude, in order to satisfy both the demand for low volume error and the ability to correctly model shortening, the stiffness coefficient must be small and the fibre count (and therefore particle count)

high. Of course, higher particle count means slower execution, so a balance between those parameters has to be found. Nevertheless, the partner facility found the method suitable for the target application.

4.3 Overall Time Performance

In order to evaluate the performance of the method, computation times of individual steps of the simulation were measured. 500 iterations of the simulation loop were set for the test. This number was chosen solely to obtain significant volume data for the time measurements. In real application, less or more iterations might be needed, depending on the desired precision. Note however, that for the tested models, 500 iterations were enough to reach a state in which the MSS start oscillating around the same values even for the lowest particle count. Only the 13 thigh muscles were used for this test.

Four different particle resolutions (i.e. numbers of particles per fibre) were used for the evaluation – 20, 40, 60 and 80. While twenty particles per fibre are sufficient to obtain acceptable output (i.e. the model does not diverge to unnatural shapes), higher resolution might offer better ratio between convergence speed and output quality. The number of fibres per muscle was always 64. While the speed of the MSS simulation depends on the total number of particles, the speed of CD depends only on the number of the boundary particles, as only those are involved in CD. The total number of particles of the muscles / the number of boundary particles for individual resolution are: 1344 / 660 for resolution 20, 2624 / 1220 for 40, 3904 / 1780 for 60 and 5184 / 2340 for 80. There were three bones significant for the used set of muscles. The same resolution (number of spheres representing the bone) was used in all cases: 8527 spheres for the pelvis, 1568 for femur (thigh) and 1709 for tibia (shank).

Table 1 shows the result of the test. The “MSS” line contains the time of simulation of the MSS (the 26-cubic lattice layout was used). “Muscle CD” line contains the time of CD of muscles vs. muscles and the “Bones CD” contains the time of muscles vs. bones CD. The times do not include any pre-processing. Lastly, because the results are also affected by the pose of the model (i.e. there might be more collisions in one position than some other), the provided times are actually an average of five different poses of the model.

The used MSS solver is only a basic, non-parallel implementation which can be improved on to achieve better results. The author of the solver, Zelený (2011), also implemented a GPU version

Table 1: Computation times in seconds for individual parts of the method after 500 simulation iterations.

Resolution Part	20	40	60	80
MSS	18.66	39.97	60.53	80.70
Muscle CD	8.15	13.50	24.00	29.53
Bones CD	30.03	37.43	56.03	67.10
TOTAL	56.84	90.9	140.56	177.33

obtaining on average 5-8x speed-up, depending on the number of particles. However, his latest implementation cannot handle large systems due to memory issues and would have to be modified first. As this work was not finished yet, the CPU solver was used here instead. Also note that the higher computation time of CD between bones and muscles than muscles vs. muscles is expected due to the higher overall resolution of the bones (see above).

Table 1 shows a linear relation between the number of particles and CD time consumption, which is rather disappointing, as the relation should be, in ideal case, logarithmic due to the hierarchical structure. On average, 76% of the CD computation time is the hierarchy traversal and 24% the piecewise testing. In the time counted as the piecewise testing there is also accounted the collision response, i.e. computing the vectors by which the spheres are moved (see section 3.3), therefore the traversal is considered the bottleneck.

Possible explanation of the slower-than-expected performance is the loose fit of AABBs. As all the objects occupy more or less the same space (they are wrapped around the same bone), the first levels of the hierarchy of two muscles will always intersect, resulting in many “dead end” traversals. However, even if for example OBBs were used instead, the situation would not improve much as there would still be significant overlap of the bounding volumes due to the nature of the data.

5 CONCLUSIONS

A mass-spring model designed for representing muscles in a musculoskeletal model was presented. The mass-points are obtained by sampling the fibres of the muscle. Three different ways of connecting the particles by springs were tested. The importance of the layout itself was found to be only marginal, while the number of springs per particle influences the convergence speed the most.

A collision detection and response mechanism

was designed, implemented and tested. The mechanism employs bounding volume hierarchies to enhance the speed, the particles of the mass-spring system are utilized as spheres of various radii approximating the surface of each object. This does not only speed up CD tests, but also bypasses the need to propagate changes of the shape between the surface model and the mass-spring model in each iteration of the simulation.

The proposed solution was evaluated by medical experts and deemed as suitable for the purpose of quick coarse simulations of moving human patient and subsequent analysis of muscle deformation.

Other methods for collision detection should be investigated further as the future work. It is the bottleneck of the solution and the current mechanism does not perform as well as expected. However, the given task is quite complex, therefore there is no guarantee that other algorithms will perform better.

ACKNOWLEDGEMENTS

This work has been supported by the Information Society Technologies Programme of the European Commission under the project VPHOP (FP7-ICT-223865), the European Regional Development Fund (ERDF) – project NTIS (New Technologies for Information Society), European Centre of Excellence, CZ.1.05/1.1.00/0.2.0090 and by the project SGS-2013-029 – Advanced Computing and Information Systems.

REFERENCES

- Faure, F., Barbier, S., Allard, J., Falipou, F., 2008. Image-based collision detection and response between arbitrary volume objects. In *Proceedings of the 2008 ACM SIGGRAPH/Eurographics Symposium on Computer Animation* (SCA '08). Eurographics Association, Aire-la-Ville, Switzerland, Switzerland, 155-162.
- Ju, T., Schaefer, S., Warren, J., 2005. Mean value coordinates for closed triangular meshes. In *ACM SIGGRAPH 2005 Papers* (SIGGRAPH '05), Markus Gross (Ed.). ACM, New York, NY, USA, 561-566.
- Lefebvre, S., Hoppe, H., 2006. Perfect spatial hashing. In *ACM SIGGRAPH 2006 Papers* (SIGGRAPH '06). ACM, New York, NY, USA, 579-588.
- Larsson, T., Akenine-Möller, T., 2006. *A dynamic bounding volume hierarchy for generalized collision detection*. *Computer Graphics* 30, 3. Pergamon Press, Inc. Elmsford, NY, USA, 450-459.
- Lee, D., Glueck, M., Khan, A., Fiume, E., Jackson, K., 2012. Modeling and Simulation of Skeletal Muscle for

- Computer Graphics: A Survey. *Foundations and Trends in Computer Graphics and Vision*, 7, 4. Now Publishers Inc. Hanover, MA, USA 229-276.
- Lee, S.-H., Sifakis, E., Terzopoulos, D., 2009. Comprehensive biomechanical modeling and simulation of the upper body. *ACM Trans. Graph.* 28, 4, Article 99, 17 pages.
- Long, J., Burns, K., Yang, J., 2011. Cloth modeling and simulation: a literature survey. In *Proceedings of the Third international conference on Digital human modeling (ICDHM'11)*, Vincent G. Duffy (Ed.). Springer-Verlag, Berlin, Heidelberg, 312-320.
- Mendoza, C., O'Sullivan, C., 2006. Interruptible collision detection for deformable objects. *Comp. Graphics* 30, 3. Pergamon Press, Inc. Elmsford, NY, USA 432-438.
- Nedel, L. P., Thalmann, D., 1998. Real Time Muscle Deformations using Mass-Spring Systems. In *Proceedings of the Computer Graphics International 1998 (CGI '98)*. IEEE Computer Society, Washington, DC, USA, 156-.
- Tang, M., Manocha, D., Tong, R., 2009. Multi-core collision detection between deformable models. In *2009 SIAM/ACM Joint Conference on Geometric and Physical Modeling (SPM '09)*. ACM, New York, NY, USA, 355-360.
- Tang, Y.-M., Hui, K.-C., 2009. Simulating tendon motion with axial mass-spring system. *Comp. Graphics* 33, 2. Pergamon Press, Inc. Elmsford, NY, USA, 162-172.
- Teschner, M., Heidelberger, B., Mueller, M., Pomeranets, D., Gross, M., 2003. Optimized Spatial Hashing for Collision Detection of Deformable Objects. In *Proceedings of Vision, Modeling, Visualization VMV'03*. Munich, Germany, 47-54.
- Villard, P.-F., Bourne, W., Bello, F., 2008. Modelling Organ Deformation Using Mass-Springs and Tensional Integrity. In *Proceedings of the 4th international symposium on Biomedical Simulation (ISBMS '08)*, Fernando Bello and P. J. Edwards (Eds.). Springer-Verlag, Berlin, Heidelberg, 221-226.
- Zelený I.: Vzájemná transformace 3D objektů reprezentovaných trojúhelníkovým povrchem [Morphing of 3D objects represented by triangular surface]. Diploma thesis, University of West Bohemia, Faculty of applied sciences, 2011. (Available only in Czech language).